

OS Project 1



Operating System, Spring 2020



Assignment 0:

Compiling Linux kernel + Adding a simple system call

1 Install Ubuntu and VirtualBox

1. Ubuntu Desktop 18.04 LTS Download

<https://ubuntu.com/#download>



2. VirtualBox Download and Install

<https://www.virtualbox.org/wiki/Downloads>



1 Install Ubuntu and VirtualBox

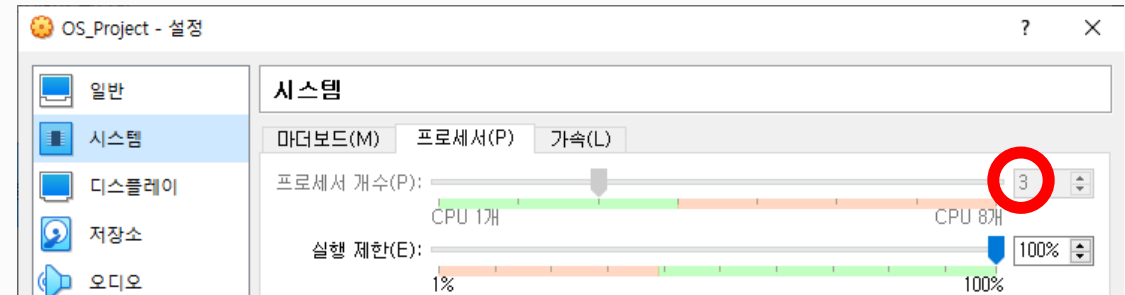
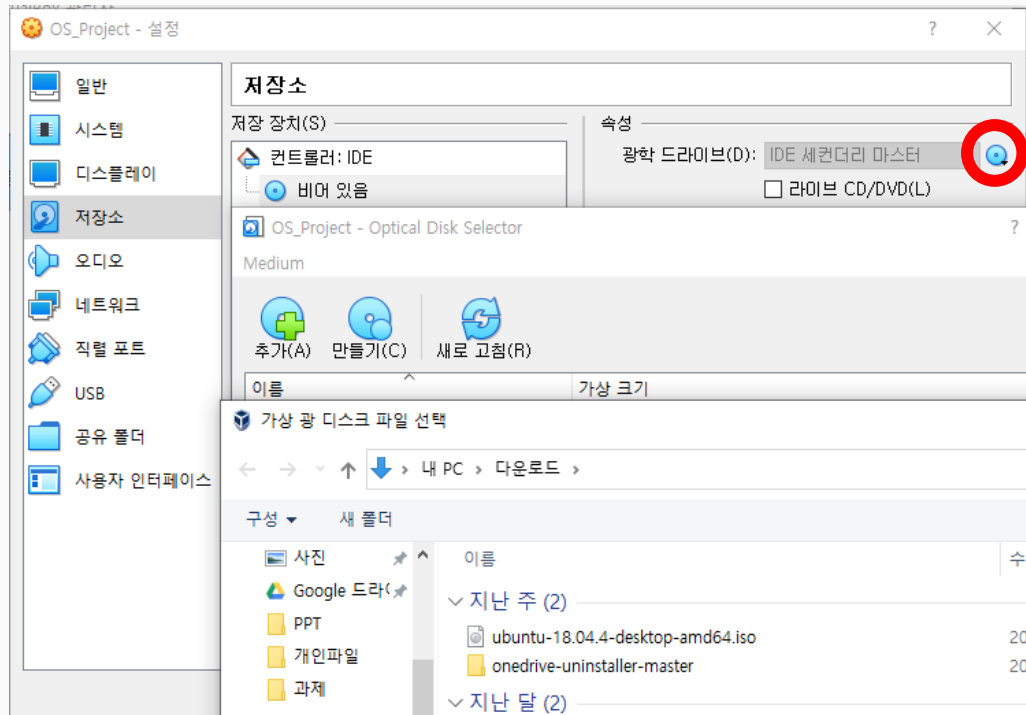
3. Virtual Box Configuration

Memory : roughly 4096MB

Virtual Hard disk(VDI) : at least 50GB

해당 가상환경의 설정에서
저장소-컨트롤러:IDE-비어있음-광학드라이브 Click
다운받은 ubuntu-18.04 선택 후 확인

Set the number of processor to be 1 less than the number of physical cores of the PC CPU



1 Install Ubuntu and VirtualBox

4. Boot Ubuntu

5. Modify grub Configuration

```
$ sudo vi /etc/default/grub
```

Adding # in the following two lines

```
GRUB_TIMEOUT_STYLE=hidden  
GRUB_TIMEOUT=0
```

```
# /boot/grub/grub.cfg.  
# For full documentation of the options in this file, see:  
#   info -f grub -n 'Simple configuration'  
  
GRUB_DEFAULT=0  
#GRUB_TIMEOUT_STYLE=hidden  
#GRUB_TIMEOUT=0  
GRUB_DISTRIBUTOR=`lsb_release -i -s 2> /dev/null || echo Debian`  
GRUB_CMDLINE_LINUX_DEFAULT="quiet splash"
```

```
$ sudo update-grub
```

```
ryotta205@ryotta205-VirtualBox:~$ sudo vi /etc/default/grub  
ryotta205@ryotta205-VirtualBox:~$ sudo update-grub  
Sourcing file `/etc/default/grub'  
grub 설정 파일을 형성합니다 ...  
Found memtest86+ image: /boot/memtest86+.elf  
Found memtest86+ image: /boot/memtest86+.bin  
완료되었습니다
```

Upon rebooting the following screen is displayed.

```
GNU GRUB 버전 2.02  
  
*Ubuntu  
Ubuntu용 고급 설정  
Memory test (memtest86+)  
Memory test (memtest86+, serial console 115200)
```

2 Install new Linux kernel

1. Check the current kernel version and download the 5.5.16 version kernel.

\$ uname -r

```
ryotta205@ryotta205-VirtualBox:~$ uname -r
5.3.0-28-generic
```

2. Download new kernel (kernel version should be 5.5.16)

\$ wget <https://cdn.kernel.org/pub/linux/kernel/v5.x/linux-5.5.16.tar.xz>

or

<https://www.kernel.org/>

mainline:	5.6	2020-03-29	[tarball]	[pgp]	[patch]	
stable:	5.6.3	2020-04-08	[tarball]	[pgp]	[patch]	[inc. patch]
stable:	5.5.16	2020-04-08	[tarball]	[pgp]	[patch]	[inc. patch]
longterm:	5.4.31	2020-04-08	[tarball]	[pgp]	[patch]	[inc. patch]
longterm:	4.19.114	2020-04-02	[tarball]	[pgp]	[patch]	[inc. patch]

3. Kernel decompression and Package download

\$ sudo mv linux-5.5.16.tar.xz /usr/src/

\$ cd /usr/src

\$ sudo xz -d linux-5.5.16.tar.xz

\$ sudo tar xf linux-5.5.16.tar

\$ cd linux-5.5.16

\$ ls

```
ryotta205@ryotta205-VirtualBox:/usr/src$ sudo tar xf linux-5.5.16.tar
ryotta205@ryotta205-VirtualBox:/usr/src$ cd linux-5.5.16
ryotta205@ryotta205-VirtualBox:/usr/src/linux-5.5.16$ ls
COPYING      Kconfig      README      crypto      init         mm           security     virt
CREDITS      LICENSES     arch        drivers     ipc          net          sound
Documentation MAINTAINERS  block       fs          kernel       samples     tools
Kbuild       Makefile     certs       include     lib          scripts     usr
ryotta205@ryotta205-VirtualBox:/usr/src/linux-5.5.16$
```

\$ sudo apt update

\$ sudo apt install build-essential libncurses5 libncurses5-dev bin86 kernel-package libssl-dev bison flex libelf-dev

2 Install new Linux kernel - Compile

4. Use the current kernel configuration as it is.

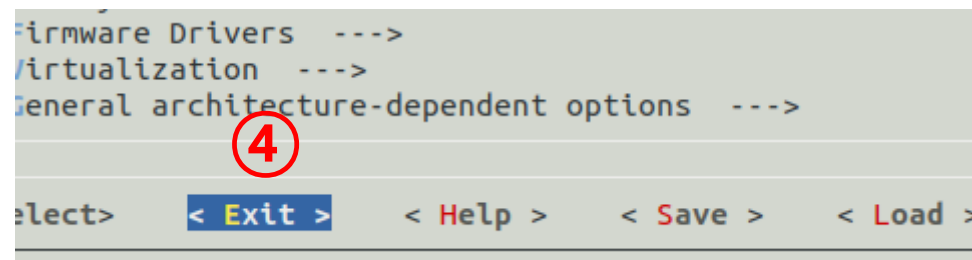
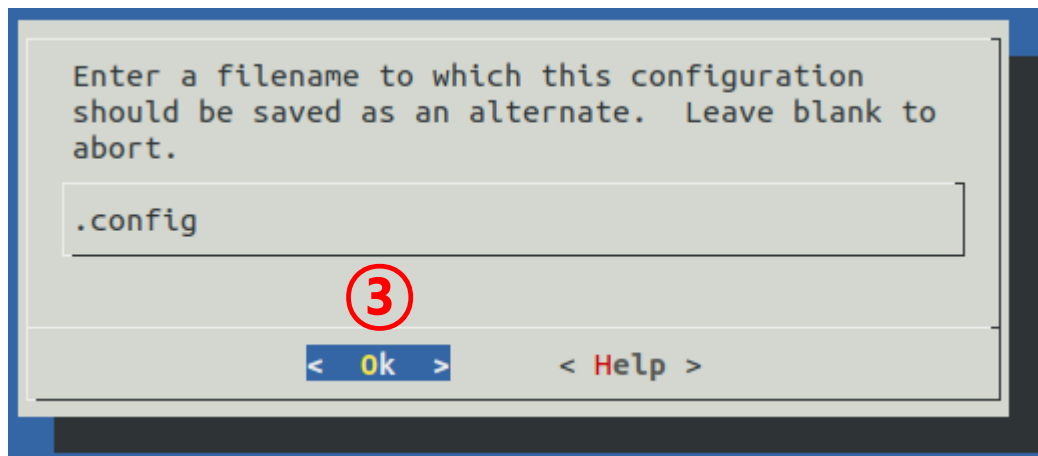
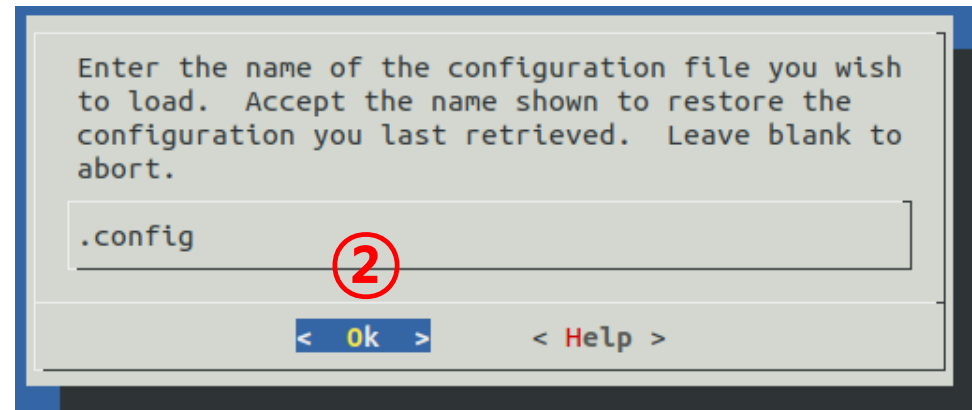
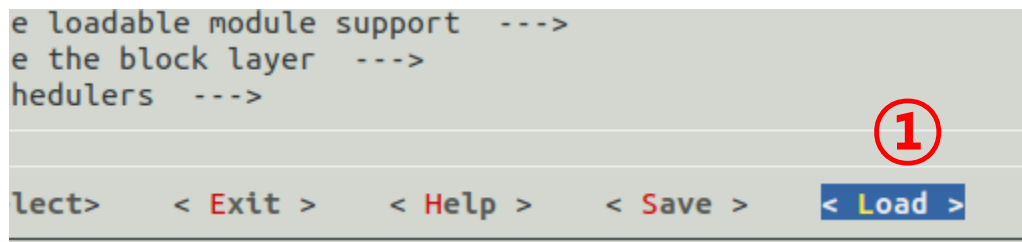
```
$ cd /usr/src/linux-5.5.16
```

```
$ sudo cp /boot/config-5.3.0-28-generic ./config
```

```
$ sudo make menuconfig
```

Using the arrow keys

Load – Ok – Save – Ok - Exit



2 Install new Linux kernel - Compile

5. Compile and Install

Compile and install the downloaded kernel and kernel module

```
$ sudo make -j3
```

Number of after j = number of allocated CPU cores

```
$ grep -c processor /proc/cpuinfo
```

 Use this to check the number of allocated CPU cores

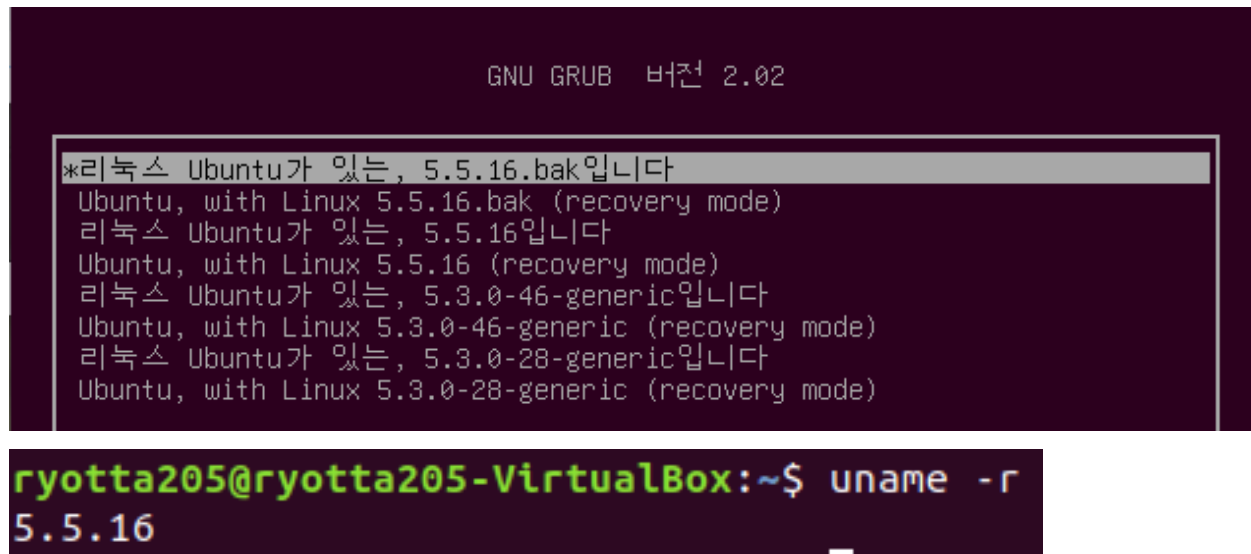
```
$ sudo make modules_install
```

```
$ sudo make install
```

```
$ sudo reboot
```

Depending on performance, 30 minutes ~ 5 hours is required

Upon rebooting the following screen is displayed. (Ubuntu용 고급 설정 선택)



```
GNU GRUB 버전 2.02

*리눅스 Ubuntu가 있는, 5.5.16.bak입니다
Ubuntu, with Linux 5.5.16.bak (recovery mode)
리눅스 Ubuntu가 있는, 5.5.16입니다
Ubuntu, with Linux 5.5.16 (recovery mode)
리눅스 Ubuntu가 있는, 5.3.0-46-generic입니다
Ubuntu, with Linux 5.3.0-46-generic (recovery mode)
리눅스 Ubuntu가 있는, 5.3.0-28-generic입니다
Ubuntu, with Linux 5.3.0-28-generic (recovery mode)

ryotta205@ryotta205-VirtualBox:~$ uname -r
5.5.16
```


3 Adding a System Call

1. Make a simple System call

```
$ cd /usr/src/linux-5.5.16  
$ sudo vi kernel/hello.c
```

hello.c

```
#include <linux/kernel.h>  
#include <linux/linkage.h>  
#include <linux/syscalls.h>  
int sys_hello(int x, int y){  
    printk("x : %d",x);  
    printk("y : %d",y);  
    printk("HELLO_SYSTEMCALL_x*y = %d\n",x*y);  
    return x*y;  
}  
  
SYSCALL_DEFINE2(hello, int, x, int, y)  
{  
    return sys_hello(x,y);  
}
```

2. Register a System call in System call table

```
$ cd /usr/src/linux-5.5.16  
$ sudo vi arch/x86/entry/syscalls/syscall_64.tbl  
436      common hello      __x64_sys_hello
```

syscall_64.tbl

```
433      common fspick      __x64_sys_fspick  
434      common pidfd_open  __x64_sys_pidfd_open  
435      common clone3      __x64_sys_clone3/ptregs  
436      common hello       __x64_sys_hello  
  
#  
# x32-specific system call numbers start at 512 to avoid cache impact
```

3 Adding a System Call

3. Declaration a System call prototype

```
$ cd /usr/src/linux-5.5.16
```

```
$ sudo vi include/linux/syscalls.h
```

Register under '`#ifndef CONFIG_ARCH_HAS_SYS_CALL_WRAPPER`'
`asmlinkage int sys_hello(int x, int y);`

```
285 * for architectures overriding the syscall calling convention, do not
286 * include the prototypes if CONFIG_ARCH_HAS_SYSCALL_WRAPPER is enabled.
287 */
288 #ifndef CONFIG_ARCH_HAS_SYSCALL_WRAPPER
289 asmlinkage int sys_hello(int x, int y);
290 asmlinkage long sys_io_setup(unsigned nr_reqs, aio_context_t __user *ctx);
291 asmlinkage long sys_io_destroy(aio_context_t ctx);
```

4. Modify Makefile

```
$ cd /usr/src/linux-5.5.16
```

```
$ sudo vi kernel/Makefile
```

Add '`hello.o`' to the end of '`obj-y`'

```
obj-y      = fork.o exec_domain.o panic.o \
             cpu.o exit.o softirq.o resource.o \
             sysctl.o sysctl_binary.o capability.o ptrace.o user.o \
             signal.o sys.o umh.o workqueue.o pid.o task_work.o \
             extable.o params.o \
             kthread.o sys_ni.o nsproxy.o \
             notifier.o ksysfs.o cred.o reboot.o \
             async.o range.o smpboot.o ucount.o hello.o
```

3 Adding a System Call

5. Recompile

Compile only kernel image and exchange (5~10 minute)

```
$ cd /usr/src/linux-5.5.16
```

```
$ sudo make bzImage -j3
```

(Number of after j = number of allocated CPU cores)

```
$ sudo cp arch/x86/boot/bzImage /boot/vmlinuz-5.5.16
```

```
$ reboot
```

```
root@ryotta205-VirtualBox:/usr/src/linux-5.5.16# make bzImage -j3
```

```
DESCEND  objtool
```

```
CALL     scripts/atomic/check-atomics.sh
```

```
CALL     scripts/checksyscalls.sh
```

```
CC       init/main.o
```

```
OBJCOPY  arch/x86/boot/setup.bin
```

```
BUILD    arch/x86/boot/bzImage
```

```
Setup is 17724 bytes (padded to 17920 bytes).
```

```
System is 9157 kB
```

```
CRC ca10ee9b
```

```
Kernel: arch/x86/boot/bzImage is ready (#9)
```

```
root@ryotta205-VirtualBox:/usr/src/linux-5.5.16# sudo cp arch/x86/boot/bzImage /boot/vmlinuz-5.5.16
```

3 Adding a System Call

6. Test

```
$ vi test.c
$ gcc test.c
$ ./a.out
$ dmesg
```

test.c

```
#include <unistd.h>
#include <stdio.h>

#define __NR_hello 436

int main(void){

    int n=0;
    n = syscall(__NR_hello, 9, 7);
    printf("return value : %d\n", n);
    return 0;
}
```

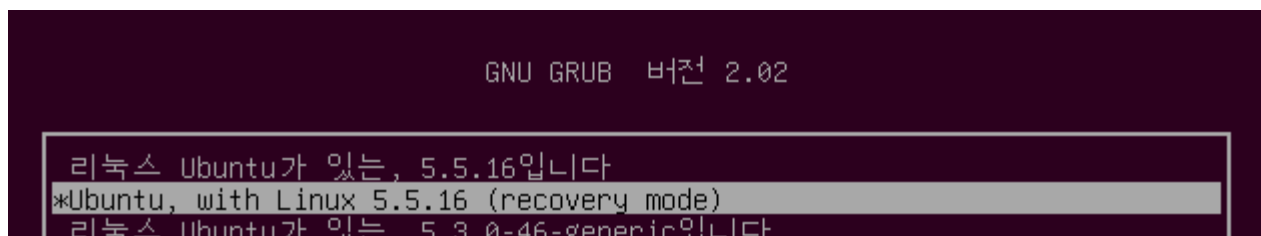
```
ryotta205@ryotta205-VirtualBox:~$ vi test.c
ryotta205@ryotta205-VirtualBox:~$ gcc test.c
test.c: In function 'main':
test.c:9:6: warning: implicit declaration of function 'syscall'; did you mean 'scanf'? [-Wimplicit-function-declaration]
    n = syscall(__NR_hello, 9, 7);
        ^~~~~~
        scanf
ryotta205@ryotta205-VirtualBox:~$ ./a.out
return value : 63
ryotta205@ryotta205-VirtualBox:~$
```

```
[ 127.075280] x : 9
[ 127.075281] y : 7
[ 127.075283] HELLO_SYSCALL_x*y = 63
ryotta205@ryotta205-VirtualBox:~$
```

4 Troubleshooting

1. 만약 부팅 중 Ubuntu 로고에서 오랫동안 멈춰 있다면 재부팅
여러 번의 재부팅 이후로도 안된다면 2번 진행
2. 만약 부팅 중 Started GNOME Display Manager에서 멈춰 있다면
Recovery mode 진입 후
root – Enter
\$ apt-get purge gdm3
\$ apt-get install gdm3 ubuntu-desktop
\$ reboot

```
Starting Hostname Service...  
[ OK ] Started Permit User Sessions.  
Starting GNOME Display Manager  
Starting Hold until boot proce  
[ OK ] Started Authorization Manager.  
[ OK ] Started Accounts Service.  
[ OK ] Started Hostname Service.  
Starting Network Manager Script  
[ OK ] Started Network Manager Script  
[ OK ] Started The PHP 7.2 FastCGI Pro  
[ OK ] Started Disk Manager.  
[ OK ] Started Clean php session files  
[ OK ] Started GNOME Display Manager.
```



```
root@ryotta205-VirtualBox:~# apt-get purge gdm3  
Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
The following packages will be REMOVED:  
gdm3 ubuntu-desktop
```

```
root@ryotta205-VirtualBox:~# apt-get install gdm3 ubuntu-desktop  
Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
Suggested packages:  
gnome-orca libpam-fprintd  
The following NEW packages will be installed:  
gdm3 ubuntu-desktop
```



Assignment 1:

Adding “sys_mygetsid()” system call

1. Add sys_mygetsid()

- **sys_mygetsid(int* buffer)**
 - 매개변수로 전달한 포인터 변수에 자신의 학번을 전달하는 system call

Test user 프로그램 소스코드: sys_mygetsid()

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

#define __NR_hello 436
#define __NR_mygetsid 437

int main(void){

    int *n;

    n = (int*)calloc(1,sizeof(int));
    printf("before buffer value : %d\n", *n);
    syscall(__NR_mygetsid, n);
    printf("after buffer value : %d\n", *n);
    return 0;
}
```

실행결과

```
ryotta205@ryotta205-VirtualBox:~$ ./a.out
before buffer value : 0
after buffer value : 2015114398
ryotta205@ryotta205-VirtualBox:~$
```

dmesg 확인결과

```
[ 1157.676032] system call : sys_mygetsid()
[ 1157.676033] system call : sys_mygetsid() : 2015114398
```

2. 제출 및 평가

제출 자료 (lms에 제출)

1. Linux kernel image (vmlinuz) file
2. source code of the system call
3. source code of the test user program
4. Result of execution (Screenshot)

평가

1. sys_hello() System call 추가 (20점)
2. System call 추가 후 Compile 성공 (10점)
3. Compile된 커널로 부팅 성공 (10점)
4. sys_mygetsid() 구현 (20점)



Assignment 2:

Adding “sys_sread()”, “sys_swrite()” system call

1. Add sys_swrite(), sys_sread()

- **sys_swrite(int fd, char* buf, int len): A secure version of sys_write() system call**
 - buf 포인터 변수가 가리키는 메모리 공간에 저장된 데이터를 암호화(encryption)한 뒤, fd가 가리키는 파일에 len 크기 만큼 저장하는 system call
 - Encryption 방법: 데이터의 모든 비트를 반전 시킴
예) 원본 데이터 (01110000) → 암호화된 데이터 (10001111)
- **sys_sread(int fd, char* buf, int len): A secure version of sys_read() system call**
 - fd가 가리키는 파일에서 len 크기만큼 읽어 복호화(decryption)한 뒤, buf포인터 변수가 가리키는 메모리 공간에 저장하는 system call
 - 파일에서 읽은 데이터는 복호화한 뒤 buf에 저장해야 함

1. Add sys_write(), sys_read()

Test user 프로그램 소스코드: sys_write()

```
#include <stdio.h>
#include <string.h>
#include <fcntl.h>
#include <unistd.h>
#include <stdlib.h>

#define __NR_hello 436
#define __NR_mygetsid 437
#define __NR_swrite 438
#define __NR_sread 439

#define BUFSIZE 100

void to_binary(char* buf){
    unsigned char p;

    for(int j = 0; j < strlen(buf) ; j++){
        p = 0x80;
        for(int i=0; i<8 ;i++){
            if(buf[j] & p){
                printf("1");
            }else{
                printf("0");
            }
            p = p >> 1;
        }
        printf(" ");
    }
    printf("\n");
}
```

```
int main()
{
    char *temp1;
    char *temp2;
    int fd;

    temp1 = (char*)calloc(100,sizeof(char));
    temp2 = (char*)calloc(100,sizeof(char));

    printf("Input data : ");
    scanf("%s",temp1);
    printf("Input binary data : ");
    to_binary(temp1);

    if ( 0 < ( fd = open( "./test.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644))) {
        syscall(__NR_swrite, fd, temp1, strlen(temp1));
        close(fd);
    }
    else{
        printf( "Open Error\n");
        exit(-1);
    }

    if ( 0 < ( fd = open( "./test.txt", O_RDONLY, 0644))) {
        read(fd, temp2, BUFSIZE);
        close(fd);
    }
    else{
        printf( "Open Error.\n");
        exit(-1);
    }

    printf("sys_read() data : %s\n",temp2);
    printf("sys_read() binary data : ");
    to_binary(temp2);

    return 0;
}
```

1. Add sys_swrite(), sys_sread()

Test user 프로그램 소스코드: sys_sread()

```
#include <stdio.h>
#include <string.h>
#include <fcntl.h>
#include <unistd.h>
#include <stdlib.h>

#define __NR_hello 436
#define __NR_mygetsid 437
#define __NR_swrite 438
#define __NR_sread 439

#define BUFSIZE 100

void to_binary(char* buf){
    unsigned char p;

    for(int j = 0; j < strlen(buf) ; j++){
        p = 0x80;
        for(int i=0; i<8 ;i++){
            if(buf[j] & p){
                printf("1");
            }else{
                printf("0");
            }
            p = p >> 1;
        }
        printf(" ");
    }
    printf("\n");
}
```

```
int main()
{
    char *temp2;
    char *temp3;
    int fd;

    temp2 = (char*)calloc(100,sizeof(char));
    temp3 = (char*)calloc(100,sizeof(char));

    if ( 0 < ( fd = open( "./test.txt", O_RDONLY, 0644))) {
        read(fd, temp2, BUFSIZE);
        close(fd);
    }
    else{
        printf( "Open Error.\n");
        exit(-1);
    }

    printf("sys_read() data : %s\n",temp2);
    printf("sys_read() binary data : ");
    to_binary(temp2);

    if ( 0 < ( fd = open( "./test.txt", O_RDONLY, 0644))) {
        syscall(__NR_sread, fd, temp3, BUFSIZE);
        close(fd);
    }
    else{
        printf( "Open Error.\n");
        exit(-1);
    }

    printf("sys_sread() data : %s\n",temp3);
    printf("sys_sread() binary data : ");
    to_binary(temp3);

    return 0;
}
```

1. Add sys_swrite(), sys_sread()

Test user 프로그램 소스코드: sys_sread() 실행결과

```
ryotta205@ryotta205-VirtualBox:~$ ./a.out
Input data : Linux
Input binary data : 01001100 01101001 01101110 01110101 01111000
sys_read() data : ♦♦♦♦♦
sys_read() binary data : 10110011 10010110 10010001 10001010 10000111
```

dmesg 확인결과

```
[ 844.154938] sys_swrite()
```

Test user 프로그램 소스코드: sys_sread() 실행결과

```
ryotta205@ryotta205-VirtualBox:~$ ./a.out
sys_read() data : ♦♦♦♦♦
sys_read() binary data : 10110011 10010110 10010001 10001010 10000111
sys_sread() data : Linux
sys_sread() binary data : 01001100 01101001 01101110 01110101 01111000
```

dmesg 확인결과

```
[ 940.868545] sys_sread()
```

2. Make a library for sys_write(), sys_read()

How to make a library?

Static Library

```
$ vi hello.c
$ gcc -c hello.c
$ ar rc libhello.a hello.o
$ vi static.c
$ gcc static.c -o static -L./ -lhello
$ ./static
```

hello.c

```
#include <unistd.h>

#define __NR_hello 436

int hello(int n, int m){
    return syscall(__NR_hello, n, m);
}
```

static.c

```
#include <unistd.h>
#include <stdio.h>

int main(){
    int n;

    n = hello(9,9);
    printf("hello return : %d\n", n);

    return 1;
}
```

결과

```
ryotta205@ryotta205-VirtualBox:~$ ./static
hello return : 81
```

2. Make a library for sys_write(), sys_read()

How to make a library?

Shared Library

```
$ vi hello.c
$ gcc -fPIC -c hello.c -o hello.o
$ gcc -shared -o libhello.so.0.0.0 hello.o
$ sudo cp libhello.so.0.0.0 /usr/lib/libhello.so.0.0.0
$ sudo ln -s /usr/lib/libhello.so.0.0.0 /usr/lib/libhello.so
$ sudo vi /etc/ld.so.conf.d/libhello.conf
$ sudo ldconfig
$ vi shared.c
$ gcc shared.c -o shared -lhello
$ ./shared
$ ldd shared
```

ldd shared

```
ryotta205@ryotta205-VirtualBox:~$ ldd shared
linux-vdso.so.1 (0x00007fffa2137000)
libhello.so => /usr/lib/libhello.so (0x00007f8baa2dc000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f8ba9eeb000)
/lib64/ld-linux-x86-64.so.2 (0x00007f8baa6e0000)
```

hello.c

```
#include <unistd.h>

#define __NR_hello 436

int hello(int n, int m){
    return syscall(__NR_hello, n, m);
}
```

shared.c

```
#include <unistd.h>
#include <stdio.h>

int main(){
    int n;

    n = hello(9,9);
    printf("hello return : %d\n", n);

    return 1;
}
```

libhello.conf

```
# libc default configuration
/usr/lib
```

결과

```
ryotta205@ryotta205-VirtualBox:~$ ./static
hello return : 81
```

2. Make a library for `sys_write()`, `sys_read()`

- **Make a library providing user-level wrapper functions for `sys_write()` and `sys_read()` system calls.**
- **Library function name should be `sread()` and `swrite()` for `sys_read()` and `sys_write()`, respectively**

3. 제출 및 평가

제출 자료 (lms에 제출)

1. Linux kernel image (vmlinuz) file
2. source code of the system call
3. source code of the test user program
4. source code of the library
5. Result of execution (Screenshot)

평가

1. sys_sread() system call 추가(5점)
2. sys_swrite() system call 추가 (5점)
3. sys_sread() 구현 (10점)
4. sys_swrite() 구현 (10점)
5. Library 구현 (10점)

Question

Email : jhyn207@gmail.com

Message : 010 - 4905 - 1749